CHMT-Q: Causal Hierarchical Market Transformer with Quantiles

A Leakage-Aware Failure Analysis of Why Better Proper Scores Do Not Imply Tradable Edge

Dhruv Sharma[1]

[1]Affiliation and correspondence details to be completed before submission

8 August 2026

Abstract

Forecasting studies in electronic markets often report predictive loss without establishing whether the learned conditional distribution contains an executable signal. We document this gap through a bounded negative-result study of CHMT-Q, a 2,089,554-parameter Causal Hierarchical Market Transformer with Quantile outputs. CHMT-Q consumes 64 minutes of bars, 120 seconds of level-1 quotes, ten-level level-2 states, and 16 data-quality features. It emits five one-minute-ahead marginal distributions for three point-of-control offsets, residual log volume, and delta fraction. The model was trained once on 18 UTC sessions of OKX BTC-USDT spot data from June 2026, selected on two sessions, and evaluated on four development-forward sessions. The final six sessions cannot serve as a confirmatory holdout because a later schema/count projection physically read their target fields, although no values were printed, retained, analyzed, or used.

On 5,724 development-forward examples, CHMT-Q reduced weighted normalized pinball loss from 0.257193 for an empirical marginal baseline to 0.251391, a 2.26% relative improvement. That aggregate result did not transfer to the intended economic use. At the five-minute horizon, total point-of-control median mean absolute error improved by only 0.012%; Pearson and Spearman associations were 0.077 and 0.075; the predicted-to-realized standard-deviation ratio was 0.151; and every nominal 10%–90% interval crossed zero. The registered uncertainty-aware policy therefore made zero trades. A post hoc median-sign long-flat policy earned 2.16% before costs but lost 68.07% under a 12-basis-point one-way cost assumption across 484 trades. Its measured gross edge was 0.445 basis points per round trip, 53.9 times smaller than modeled costs.

The audit identifies task-decision misalignment, cost dominance, weak and shrunk conditional medians, severe dependence from overlapping windows, narrow validation, target redundancy, saturated delta tails, persistent gradient clipping, and weak baselines. These observations support a methodological conclusion rather than a profitability claim: a better proper score is insufficient evidence of tradable edge. We derive the model, objective, diagnostics, and cost arithmetic in detail, provide a falsifiable corrective protocol, and treat a cost-aware Transformer–Double-DQN architecture only as future work.

Keywords: *limit order book*, *probabilistic forecasting*, *quantile regression*, *Transformer*, *market microstructure*, *backtest overfitting*, *negative results*, *offline reinforcement learning*

1 Introduction

Modern limit-order-book models can optimize a statistically valid forecasting objective while remaining economically unusable. A proper scoring rule rewards a forecast for representing the evaluated target distribution [16]; a trading decision instead depends on an executable, conditional return distribution net of spread, fees, slippage, latency, queue effects, impact, and risk. A lower forecasting loss is therefore neither a theorem nor, by itself, reliable empirical evidence of positive net profit.

This paper studies that distinction with CHMT-Q, the Causal Hierarchical Market Transformer with Quantile outputs. Here, causal means that features, attention paths, normalization, and labels obey temporal availability. It does not denote causal inference or an identified treatment effect. CHMT-Q preserves the geometry of a ten-level order book, combines it with level-1 and bar histories, and emits noncrossing quantiles at five horizons. The design draws on sequence attention [1, 2], time-series patching and multi-horizon forecasting [5–8], and hierarchical treatment of order-book state [9–15, 18].

The experiment appears favorable if summarized by one number: aggregate weighted normalized pinball loss was 2.26% lower than an empirical marginal baseline. A deeper audit reverses the practical interpretation. Conditional medians were strongly shrunk, the target was not executable return, uncertainty intervals admitted no trades, and forced median-sign policies turned over far too often for their small gross edge. Consecutive examples also shared

63 of 64 bar observations and four of five longest-horizon label minutes, so the nominal row count substantially overstated independent information.

We ask four bounded research questions:

1. **Distributional score:** Does CHMT-Q improve its registered probabilistic forecasting metric on a temporally later development partition?
2. **Conditional signal:** Does that improvement correspond to useful median, rank, directional, or uncertainty information at five minutes?
3. **Economic feasibility:** Can a fixed decision rule earn a gross edge plausibly larger than explicit transaction costs?
4. **Failure mechanism:** Which data, target, optimization, validation, and execution factors explain the gap without unsupported causal attribution?

The contributions are deliberately narrow: a fully specified multimodal quantile Transformer and immutable experiment ledger; a formal separation of proper-score, conditional-signal, and economic claims; an audit showing how a positive aggregate score can coexist with zero uncertainty-authorized trades and negative cost-adjusted returns; a disclosure that retires a nominal holdout after process-level access; and a registered corrective design whose reinforcement-learning component remains unimplemented and unvalidated.

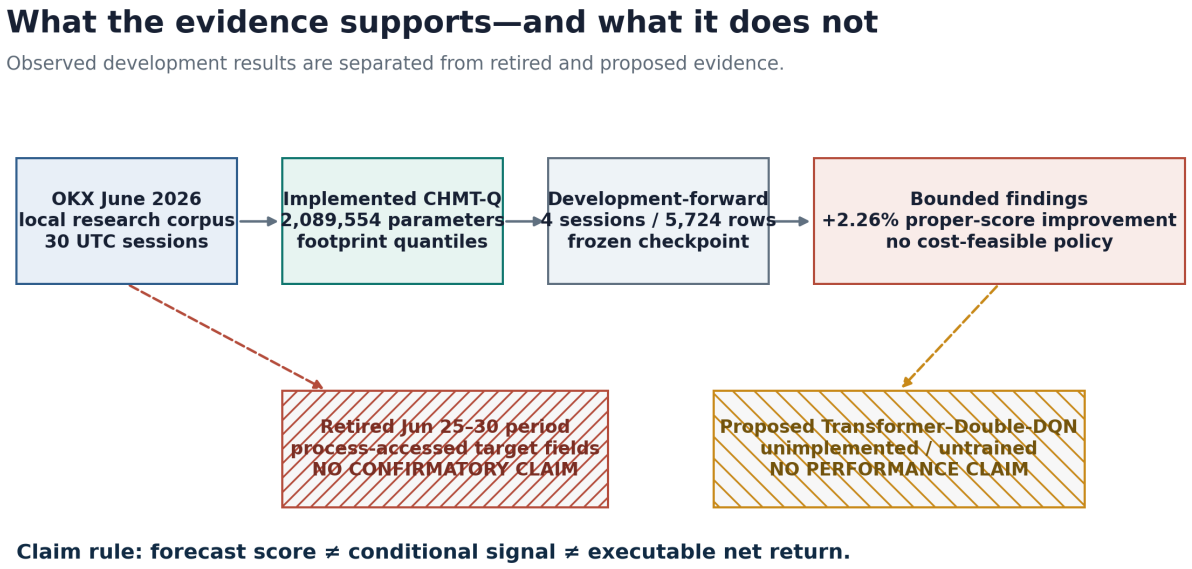


Figure 1: Evidence and claim boundary. Only frozen development artifacts support the reported findings. The process-accessed retired period and the proposed reinforcement-learning study support no performance claim.

2 Related Work

2.1 Order-book representation and multi-horizon prediction

DeepLOB established a convolutional-recurrent baseline for extracting spatial and temporal structure from limit-order-book tensors [9]. Subsequent studies used attention and Transformers for order books [10, 11], standardized prediction benchmarks [12], linked order-flow imbalance to multi-horizon return predictability [13], and showed that market, label, horizon, and evaluation choices materially change apparent performance [14, 15]. Universal price-formation evidence suggests that some representations transfer across instruments [18], but it does not establish that a particular target or execution rule transfers.

CHMT-Q follows the representational lesson that book rank, side, and time should not be flattened indiscriminately. It differs in objective: it forecasts footprint primitives rather than a direct return or trade action. That choice is central to the result because representation quality cannot repair weak alignment between the supervised target and the downstream decision.

2.2 Probabilistic forecasting

Pinball loss is a consistent score for conditional quantiles [16]. Noncrossing parameterizations can impose ordered quantiles, while conformal methods can calibrate intervals under explicit exchangeability or adaptation assumptions

[17]. CHMT-Q uses an ordered head but makes no distribution-free coverage claim. Its reported 10%–90% coverage is descriptive; no conformal correction was fit on the four-day forward partition.

The mixed record of complex time-series models against linear baselines [8] makes strong simple comparators essential. Only empirical-marginal and zero-residual baselines were registered in this run. The absence of ridge, boosted-tree, convolutional, and recurrent competitors precludes an architectural superiority claim.

2.3 Backtest validity and reinforcement learning

Financial model selection is vulnerable to repeated testing, dependence, nonstationarity, and costs. Reality checks, multiple-testing corrections, the stationary bootstrap, and backtest-overfitting diagnostics address parts of the problem [19–21, 29]. Gap-separated validation is relevant when observations share features or labels [30]. Four forward sessions are too few for stable daily-block inference, so this paper reports descriptive estimates rather than confirmatory tail probabilities.

Deep Q-learning [22], Double DQN [23], dueling decomposition [24], distributional Q-learning [25], and conservative offline objectives [26] provide tools for a future decision model. Deep reinforcement learning is high-variance [28], and offline learning adds extrapolation risk when candidate actions lack data support [27]. These methods cannot manufacture missing information or remove transaction costs.

3 Formal Problem and Claim Taxonomy

3.1 Causal information set

Definition 1 (Availability-causal forecast)**.** Let c be a decision cutoff and let every observation x_i have an availability time t_i^{avail}. The legal information set is

$$\mathcal{F}_c = \sigma\big(x_i : t_i^{\text{avail}} \le c\big)\,. \tag{3.1}$$

A forecast is availability-causal if

$$\widehat{\mathbf{Y}}_{c,1:H} = f_\theta(\mathcal{F}_c, \mathbf{M}_c)\,, \qquad H = 5, \tag{3.2}$$

where $\mathbf{M}_c$ contains only fixed configuration, masks, and normalization state permitted at cutoff c.

This definition concerns temporal availability, not causal identification. The archives lack local receive timestamps and native sequence/checksum evidence, so the implemented claim is bounded to archive event order and the documented compiler.

3.2 Forecast intervals and target transformation

For cutoff c_i and one-minute interval Δ, define

$$I_{i,h} = [c_i + (h-1)\Delta, c_i + h\Delta), \qquad h = 1, \dots, 5. \tag{3.3}$$

Let P_h, P_h^b, P_h^s denote total, buy, and sell point-of-control prices; let V_h be eligible traded volume; let D_h be signed delta; let p_{ref} be the causal reference price; let δ_p be tick size; and let ℓ_{ref} be a causal trailing log-volume reference. The five normalized targets are

$$\begin{aligned}
y_{h,1} &= \frac{P_h - p_{\text{ref}}}{\delta_p}, & y_{h,2} &= \frac{P_h^b - p_{\text{ref}}}{\delta_p},\\
y_{h,3} &= \frac{P_h^s - p_{\text{ref}}}{\delta_p}, & y_{h,4} &= \log(1+V_h) - \ell_{\text{ref}},\\
y_{h,5} &= \operatorname{clip}\left(\frac{D_h}{V_h}, -1, 1\right).
\end{aligned} \tag{3.4}$$

Each (h,k) has an independent validity mask. Missing side-specific POC values are undefined, not numerical zeros. These variables describe future footprint structure; none is an executable return.

3.3 Three evidentiary levels

Proposition 1 (A better target score need not create decision value)**.** *Let Y be the supervised target, R an executable return, and X the available state. If a model improves a proper score for $\mathcal{L}(Y \mid X)$, but $R \perp Y \mid X$ and the policy-relevant part of $\mathcal{L}(R \mid X)$ is unchanged, the improvement creates no necessary increase in optimal trading value.*

Claim level	Registered question	Evidence	Defensible conclusion
Forecast distribution	Is weighted normalized pinball below the empirical marginal baseline?	5,724 development-forward examples	Yes, by 2.26% on this partition.
Conditional signal	Are five-minute medians and intervals sufficiently selective?	MAE, correlations, sign metrics, dispersion, interval width	Weak location signal; every interval crosses zero.
Economic value	Does a frozen rule retain positive return after costs?	Frictionless and fixed-bps simulations	No under the tested assumptions.
Generalization	Does performance persist across regimes, seeds, venues, or instruments?	One month, one instrument, one venue, one seed, one fold	Not established.
Deployment	Is the model executable in a live terminal?	Offline model and replay artifacts	Not established; no live integration.

Table 1: Claim taxonomy. Evidence at one level does not automatically support the next.

Proof. For any action a and cost $C(a, X)$, conditional decision value is

$$V(a \mid X) = \mathbb{E}[aR - C(a, X) \mid X]. \tag{3.5}$$

The score for Y does not occur in this expression. Conditional independence gives $\mathbb{E}[R \mid X, Y] = \mathbb{E}[R \mid X]$, so a more accurate representation of $\mathcal{L}(Y \mid X)$ leaves every $V(a \mid X)$ unchanged unless it also changes information about R, the action map, or costs. As a concrete counterexample, take $Y = X + \epsilon$ and $R = \eta$, where ϵ, η are independent, centered noises. A model using X beats a marginal forecast for Y, while every nonzero trade has zero expected gross return and negative expected net return under positive costs. □

4 Data, Partitions, and Integrity Boundary

4.1 Source and provenance

The source corpus comprises official OKX historical archives for BTC-USDT spot for June 2026 UTC: 31 level-2 files and 31 trade files containing 149,390,769 raw level-2 records and 16,918,630 trades. Compilation produced 42,930 examples across 30 UTC sessions in 168 Arrow shards (1.612 GiB). The immutable dataset identifier is `565c1d8e912ee426a1a92572d5be481ec5168d54d1ae53bd77fdd70d888e4df6`.

The corpus lacks local receive timestamps, native sequence numbers, and checksums; June-effective instrument metadata was not independently verified; and the license/provenance review was incomplete. The data owner authorized local, non-promotable research subject to those limitations. The paper makes no claim about live receive-time causality, packet-loss reconstruction, or redistribution rights.

4.2 Chronological split and retired period

Partition	Sessions	Examples	Role
Training	18	25,758	Parameter fitting
Selection	2	2,862	Checkpoint selection
Development-forward	4	5,724	Reported evaluation
Retired final period	6	8,586	No confirmatory use

Table 2: Frozen chronological partition.

The final 8,586 examples were not scored in the historical study. Later, a schema/count projection traversed all 42,930 rows and physically read projected target fields in those shards. No values were printed, retained, analyzed, or used. Nevertheless, physical access violates a strict holdout-control policy. The period is therefore not described as sealed, untouched, or confirmatory. A future study requires a newly acquired, access-controlled holdout.

4.3 Evaluation protocol

The selected checkpoint came from epoch 10, optimizer step 2,010. Predictions were made once per minute. Economic decisions were sampled every five minutes; the proxy entered at the next minute open and exited at the close

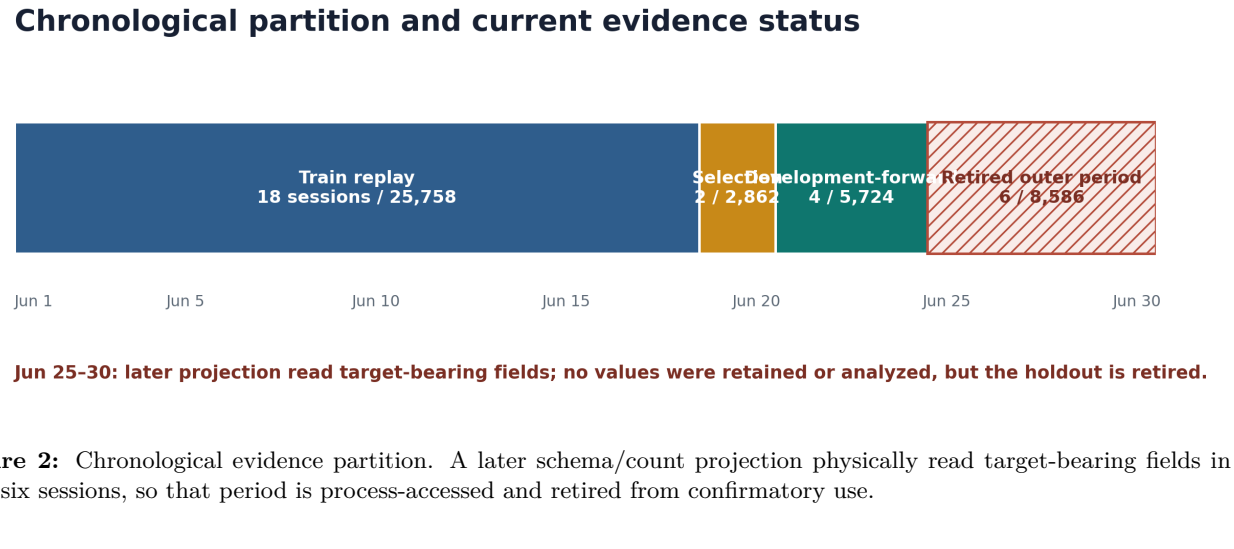


Figure 2: Chronological evidence partition. A later schema/count projection physically read target-bearing fields in the final six sessions, so that period is process-accessed and retired from confirmatory use.

four minutes later. Notional was fixed; leverage, pyramiding, endogenous impact, queue position, and partial fills were excluded. Cost scenarios were 0, 9, 12, or 20 basis points per side, with 12 bps designated as the base case.

The registered forecast endpoint was weighted normalized pinball loss against an empirical marginal forecast fit only on permitted data. The registered economic rule was interval-based long-flat selection. Median-sign long-flat and long-short rules were post hoc diagnostics. A 10,000-replicate daily-block bootstrap was computed, but four forward days are too few for stable confirmatory inference.

5 CHMT-Q Architecture

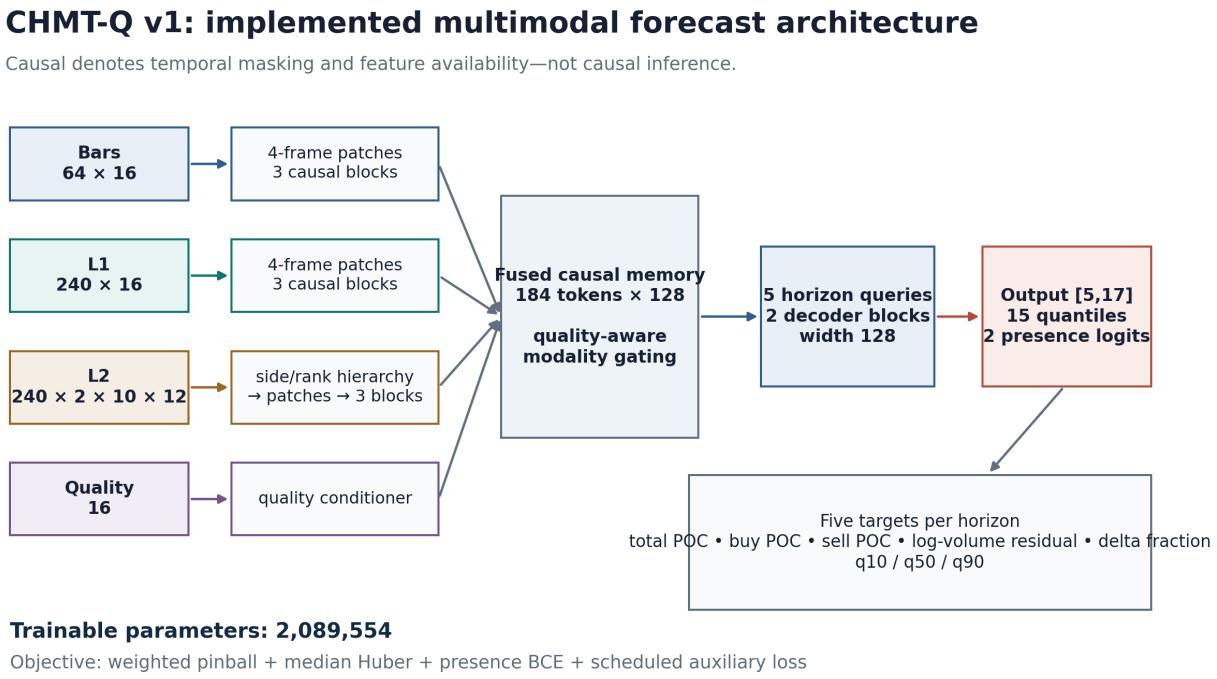


Figure 3: Implemented CHMT-Q architecture. Book rank is encoded before time; bar, L1, and L2 memories are fused with quality context and decoded by five horizon queries.

5.1 Tensor contract and shape trace

For batch size B, inputs are

$$\mathbf{X}^b \in \mathbb{R}^{B\times 64\times 16}, \quad \mathbf{X}^1 \in \mathbb{R}^{B\times 240\times 16}, \quad \mathbf{X}^2 \in \mathbb{R}^{B\times 240\times 2\times 10\times 12}, \quad \mathbf{q} \in \mathbb{R}^{B\times 16}. \tag{5.1}$$

Level-1 and level-2 span 120 seconds at nominal half-second resolution. Four microstructure frames are pooled into one two-second patch, yielding 60 tokens for each fast modality. Bars retain 64 one-minute tokens. The fused memory has $64+60+60=184$ tokens of width 128. Five learned horizon queries produce a raw tensor in $\mathbb{R}^{B\times5\times17}$: 15 values parameterize three quantiles for five targets and two logits represent buy/sell presence.

Component	**External / internal shape**	**Implemented setting**
Bars	$64\times16\to64\times128$	3 causal blocks
Level 1	$240\times16\to60\times128$	4-frame patch; 3 causal blocks
Level 2	$240\times2\times10\times12\to240\times20\times32\to60\times128$	1 rank-side block; 3 temporal blocks
Quality	$16\to128$	Three hard-aware modality gates
Main attention	width 128, 4 heads	head width 32; FFN width 256
Level attention	width 32, 4 heads	head width 8; FFN width 64
Decoder	5×128	2 blocks; 184-token cross-attention
Output	5×17	5×3 quantiles + 2 logits per horizon
Parameters	scalar total	2,089,554 trainable

Table 3: Implemented tensor and architecture contract.

5.2 Causal robust scaling and physical-time encoding

For feature j, the training-only robust scale and transformed value are

$$s_j=\max\bigl(1.4826\ \mathrm{MAD}_{\mathrm{train},j},10^{-6}\bigr),\qquad z_j=\mathrm{clip}\left(\frac{x_j-\mu_j}{s_j},-8,8\right). \tag{5.2}$$

Invalid rows are zeroed after scaling; masks remain authoritative.

For nonnegative age a, maximum age $a_{\max}$, and fixed scales $\{\tau_k\}_{k=1}^{K}$, define

$$\phi(a)=\left[\frac{a}{a_{\max}},\frac{\log(1+a)}{\log(1+a_{\max})},\sin\frac{a}{\tau_1},\cos\frac{a}{\tau_1},\dots,\sin\frac{a}{\tau_K},\cos\frac{a}{\tau_K}\right], \tag{5.3}$$

and

$$e_t(a)=W_2\,\mathrm{SiLU}(W_1\phi(a)+b_1)+b_2. \tag{5.4}$$

Bars use scales 60 through 3,840 seconds; microstructure uses scales 0.5 through 64 seconds. L2 tokens also receive learned side and rank embeddings:

$$x^{(0)}_{t,s,r}=W_{\mathrm{level}}z_{t,s,r}+e^{\mathrm{side}}_s+e^{\mathrm{rank}}_r. \tag{5.5}$$

Temporal attention head j uses a learned nonnegative recency coefficient:

$$b^{(j)}_{i\ell}=-\,\mathrm{softplus}(\gamma_j)\log\left(1+\frac{t_i-t_\ell}{s_m}\right),\qquad t_\ell\le t_i, \tag{5.6}$$

where $s_m=60$ seconds for bars and $s_m=2$ seconds for patched microstructure. This is a learned bias, not a hard decay.

5.3 Custom attention block

For width d, gain g, and $\epsilon>0$, RMSNorm [3] is

$$\mathrm{RMSNorm}(x)=g\odot\frac{x}{\sqrt{d^{-1}\sum_{k=1}^{d}x_k^2+\epsilon}}. \tag{5.7}$$

For head j, normalized states U, and memory M,

$$Q_j=UW_j^Q,\qquad K_j=MW_j^K,\qquad V_j=MW_j^V. \tag{5.8}$$

Each query and key vector is root-mean-square normalized:

$$\bar q=\frac{q}{\sqrt{d_h^{-1}\|q\|_2^2+\epsilon}},\qquad \bar k=\frac{k}{\sqrt{d_h^{-1}\|k\|_2^2+\epsilon}}. \tag{5.9}$$

With learned log-temperature θ_j, optional time bias $b_{i\ell}^{(j)}$, modality log-gate bias $g^{(m(\ell))}$, and eligibility mask $\mathcal{M}_{i\ell}$, the score is

$$s_{i\ell}^{(j)} = e^{\theta_j} \frac{\bar{q}_i^{(j)\top} \bar{k}_\ell^{(j)}}{\sqrt{d_h}} + b_{i\ell}^{(j)} + g^{(m(\ell))} + \mathcal{M}_{i\ell}. \tag{5.10}$$

For eligibility indicator $m_{i\ell}$, let $a_i = \bigvee_\ell m_{i\ell}$ and $r_{i\ell} = m_{i\ell} \exp(s_{i\ell} - s_i^{\max})$, with the maximum taken only over eligible keys. Safe masked softmax is

$$p_{i\ell} = \begin{cases} r_{i\ell} / \sum_k r_{ik}, & a_i = 1, \\ 0, & a_i = 0. \end{cases} \tag{5.11}$$

This explicit zero branch prevents NaN output when every key is masked. Multi-head output and the SwiGLU feed-forward sublayer [4] are

$$O = \mathrm{Concat}_j(P_j V_j) W^O + b^O, \tag{5.12}$$

$$\mathrm{FFN}(x) = W_o[\mathrm{SiLU}(W_g x + b_g) \odot (W_u x + b_u)] + b_o. \tag{5.13}$$

Pre-norm LayerScale residual updates are

$$\begin{aligned} u &= x + \lambda_a \odot \mathrm{Dropout}(\mathrm{Attention}(\mathrm{RMSNorm}(x))), \\ z &= u + \lambda_f \odot \mathrm{Dropout}(\mathrm{FFN}(\mathrm{RMSNorm}(u))). \end{aligned} \tag{5.14}$$

Invalid query rows are zeroed after the update.

5.4 Causal patching and hierarchical level encoding

For four projected L1 or L2 frame states x_j with masks m_j, the learned patch pool is

$$\begin{aligned} r_j &= w_2^\top \tanh(W_1 x_j + b_1) + b_2, \qquad j = 0, \ldots, 3, \\ \pi_j &= \mathrm{SafeSoftmax}_j(r_j; m_j), \qquad p = \sum_{j=0}^{3} \pi_j x_j. \end{aligned} \tag{5.15}$$

An all-invalid patch returns zero and an invalid mask.

Within one L2 snapshot, 20 bid/ask rank positions coexist at the same time, so the level block is masked but not temporally triangular. A learned book query pools level states:

$$\begin{aligned} a_{s,r} &= \mathrm{SafeSoftmax}_{s,r}\left(\frac{q_{\mathrm{book}}^\top \mathrm{RMSNorm}(x_{s,r})}{\sqrt{32}}; m_{s,r} \right), \\ d_t &= W_d \sum_{s,r} a_{s,r} x_{t,s,r} + b_d \in \mathbb{R}^{128}. \end{aligned} \tag{5.16}$$

The resulting 240 depth-frame tokens are patched to 60 and processed by three causal temporal blocks. This rank-before-time hierarchy separates same-snapshot book geometry from temporal evolution.

5.5 Quality-aware fusion and direct decoder

Let s_b, s_1, s_2 be the latest valid summaries, c_q the quality-MLP output, and $a_m \in \{0, 1\}$ a hard modality-availability indicator. For each modality,

$$\begin{aligned} \widetilde{g}_m &= a_m \, \mathrm{sigmoid}\big(w_{m,2}^\top \, \mathrm{SiLU}(W_{m,1}[s_m; c_q] + b_{m,1}) + b_{m,2}\big), \\ \alpha_m &= \frac{\widetilde{g}_m}{\sum_n \widetilde{g}_n + \epsilon}, \qquad c = \mathrm{RMSNorm}\left(c_q + \sum_m \alpha_m s_m \right). \end{aligned} \tag{5.17}$$

Encoded histories concatenate to

$$M = [M_b; M_1; M_2] \in \mathbb{R}^{B \times 184 \times 128}. \tag{5.18}$$

The decoder receives $\log(\alpha_m + \epsilon)$ as a key-modality score bias; hard-masked modalities cannot be revived by a learned gate.

For horizon $h \in \{1, \ldots, 5\}$ and forecast interval $\Delta = 60$ seconds, the initial query is

$$q_h^{(0)} = e_h + W_c \, c + \mathrm{HorizonMLP}[h/5, \log(1 + h), h\Delta]. \tag{5.19}$$

Two decoder blocks perform horizon self-attention, cross-attention to M, and SwiGLU updates. Horizon self-attention is noncausal because all five queries are formed at the same decision cutoff. No predicted horizon is fed into another.

5.6 Ordered quantile head

For each horizon-target pair, the head produces a center c and gap logits a, b:

$$\widehat{Q}_{0.50} = c, \qquad \widehat{Q}_{0.10} = c - \mathrm{softplus}(a), \qquad \widehat{Q}_{0.90} = c + \mathrm{softplus}(b). \tag{5.20}$$

Because $\mathrm{softplus}(x) = \log(1 + e^x) > 0$, the outputs satisfy $\widehat{Q}_{0.10} \le \widehat{Q}_{0.50} \le \widehat{Q}_{0.90}$. For delta fraction, the monotone map tanh preserves this order while restricting each marginal to $[-1, 1]$. Applying a saturating bounded map after gap construction can, however, collapse tail resolution; empirically, the lower delta quantile was near -1 in all forward cases and the upper quantile near $+1$ in 78.0%.

Physical reconstruction is

$$\begin{aligned} \widehat{P} &= p_{\mathrm{ref}} + \delta_p \widehat{y}_{\mathrm{POC}}, \\ \widehat{V} &= \max(0, \exp(\ell_{\mathrm{ref}} + \widehat{y}_V) - 1)\,, \\ \widehat{f}_D &= \tanh(\widehat{y}_D). \end{aligned} \tag{5.21}$$

The five-horizon output comprises marginal distributions, not a coherent joint sample path. Products of separately predicted volume and delta-fraction quantiles are not valid raw-delta quantiles.

6 Training Objective and Optimization

6.1 Pinball loss and conditional quantiles

For level $\tau \in (0, 1)$, residual $u = y - q$, and forecast q,

$$\rho_\tau(u) = u\,(\tau - \mathbb{I}\{u < 0\}) = \max\{\tau u, (\tau - 1)u\}. \tag{6.1}$$

Why pinball loss targets a conditional quantile

For a conditional distribution $F_{Y|X}$, define $R(q \mid X) = \mathbb{E}[\rho_\tau(Y - q) \mid X]$. Away from atoms,

$$\begin{aligned} \frac{\partial R}{\partial q} &= (1 - \tau)\Pr(Y < q \mid X) - \tau \Pr(Y > q \mid X) \\ &= F_{Y|X}(q) - \tau. \end{aligned} \tag{6.2}$$

Thus a minimum satisfies $F_{Y|X}(q) = \tau$. At atoms, the subgradient condition is $F(q^-) \le \tau \le F(q)$, the standard quantile definition. This consistency concerns Y; it does not convert Y into an executable return.

With horizon weights $w_h = h^{-1/2}$, target weights $v = (1, 1, 1, 0.5, 0.5)$, quantiles $\mathcal{T} = \{0.1, 0.5, 0.9\}$, and validity mask m_{bhk},

$$\mathcal{L}_Q = \frac{\sum_{b,h,k,\tau} m_{bhk} w_h v_k \rho_\tau(y_{bhk} - \widehat{Q}_{bhk\tau})}{\sum_{b,h,k,\tau} m_{bhk} w_h v_k + \epsilon}. \tag{6.3}$$

The longest-horizon total-POC contribution has nominal objective share

$$\frac{5^{-1/2} \cdot 1}{\left(\sum_{h=1}^{5} h^{-1/2}\right)\left(\sum_{k=1}^{5} v_k\right)} = 0.034596, \tag{6.4}$$

after the common quantile multiplicity cancels. Only 3.46% of the horizon-target weighting therefore directly concerns H5 total POC.

6.2 Median, presence, and auxiliary terms

For residual r and $\delta = 1$, Huber loss is

$$H_\delta(r) = \begin{cases} \frac{1}{2} r^2, & |r| \le \delta, \\ \delta(|r| - \frac{1}{2}\delta), & |r| > \delta. \end{cases} \tag{6.5}$$

The normalized median term is

$$\mathcal{L}_H = \frac{\sum_{b,h,k} m_{bhk} w_h v_k H_\delta(y_{bhk} - \widehat{Q}_{bhk,0.5})}{\sum_{b,h,k} m_{bhk} w_h v_k + \epsilon}. \tag{6.6}$$

For presence label $z \in \{0,1\}$, logit η, and presence-validity mask m^P, stable binary cross-entropy is

$$\mathrm{BCE}(\eta, z) = \max(\eta, 0) - \eta z + \log(1 + e^{-|\eta|}),$$
$$\mathcal{L}_P = \frac{\sum m^P \,\mathrm{BCE}(\eta, z)}{\sum m^P + \epsilon}. \tag{6.7}$$

The final-decoder base objective is

$$\mathcal{L}_{\text{base}} = \mathcal{L}_Q + 0.25\mathcal{L}_H + 0.20\mathcal{L}_P. \tag{6.8}$$

The same output head applied after decoder block 1 supplies a deep-supervision term. With $u \in [0,1]$ the fraction of planned optimizer steps,

$$\lambda_{\text{aux}}(u) = \begin{cases} 0.25, & u \le 0.60, \\ 0.25 - \dfrac{0.20}{0.30}(u - 0.60), & 0.60 < u < 0.90, \\ 0.05, & u \ge 0.90, \end{cases} \tag{6.9}$$

and

$$\mathcal{L} = \mathcal{L}_{\text{base}}^{\text{final}} + \lambda_{\text{aux}}(u)\mathcal{L}_{\text{base}}^{\text{aux}}. \tag{6.10}$$

Early stopping occurred after 36% of planned steps, so $\lambda_{\text{aux}} = 0.25$ throughout the run.

6.3 Optimization

The run used seed 17, AdamW with learning rate 3×10^{-4}, minimum learning rate 3×10^{-5}, $(\beta_1, \beta_2) = (0.9, 0.95)$, $\epsilon = 10^{-8}$, weight decay 0.05, and 5% warm-up. A microbatch of 32 with four-step accumulation gave an effective batch of 128. Training was deterministic FP32 on one CUDA device. It stopped at epoch 18 after 3,618 optimizer steps and selected epoch 10 at step 2,010. If $\mathcal{S}$ is the set of epoch checkpoints, the frozen selection rule was

$$\widehat{s} = \arg\min_{s \in \mathcal{S}} \mathcal{L}_{\text{validation}}(\theta_s). \tag{6.11}$$

Global norm clipping with threshold $c = 1$ maps raw gradient g to

$$\widetilde{g} = g \min\left(1, \frac{c}{\|g\|_2 + \epsilon}\right). \tag{6.12}$$

Every recorded pre-clipping norm exceeded 1.0; the median was 19.58, the 90th percentile 38.30, and the maximum 109.07. This proves operation in a persistently clipped regime, but not that clipping caused the economic failure.

Training item	**Value**
Training / selection / forward examples	25,758 / 2,862 / 5,724
Seed / folds	1 / 1
Selected checkpoint	epoch 10, step 2,010
Final step	3,618
Selected validation loss	0.210765
Empirical validation baseline	0.222446
Weight artifact	8.071 MiB
Resumable checkpoint	24.284 MiB

Table 4: Training and selection ledger. Validation improvement is selection evidence, not an independent performance estimate.

7 Evaluation Mathematics

7.1 Forecast baseline, score, and diagnostics

For permitted pre-forward observations in horizon-target cell (h, k), the empirical marginal baseline is

$$\widehat{Q}_{hk\tau}^{\text{emp}} = \inf\{q : \widehat{F}_{hk}^{\text{train}}(q) \ge \tau\}. \tag{7.1}$$

It is applied unchanged to every development-forward row. If $\mathcal{L}_{\text{model}}$ and $\mathcal{L}_{\text{base}}$ are Equation (6.3), relative improvement is

$$\Delta_{\%} = 100\frac{\mathcal{L}_{\text{base}} - \mathcal{L}_{\text{model}}}{\mathcal{L}_{\text{base}}}. \tag{7.2}$$

For H5 total POC, location diagnostics are

$$\begin{aligned} \text{MAE} &= \frac{\sum_i m_i|y_i - \widehat{q}_i|}{\sum_i m_i}, & \text{Bias} &= \frac{\sum_i m_i(\widehat{q}_i - y_i)}{\sum_i m_i}, \\ r_P &= \frac{\sum_i (y_i - \bar{y})(\widehat{q}_i - \bar{q})}{\sqrt{\sum_i (y_i - \bar{y})^2 \sum_i (\widehat{q}_i - \bar{q})^2}}, & r_S &= r_P(\text{rank}(y), \text{rank}(\widehat{q})), \\ \kappa_\sigma &= \frac{\text{sd}(\widehat{q})}{\text{sd}(y)}. \end{aligned} \tag{7.3}$$

For binary sign labels and predictions, Matthews correlation is

$$\text{MCC} = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}. \tag{7.4}$$

For interval $[L_i, U_i]$, empirical coverage, mean width, and zero-crossing rate are

$$\begin{aligned} \widehat{C} &= \frac{\sum_i m_i \mathbb{I}\{L_i \le y_i \le U_i\}}{\sum_i m_i}, & \widehat{W} &= \frac{\sum_i m_i (U_i - L_i)}{\sum_i m_i}, \\ \widehat{Z} &= \frac{\sum_i m_i \mathbb{I}\{L_i \le 0 \le U_i\}}{\sum_i m_i}. \end{aligned} \tag{7.5}$$

7.2 Decision rules and return accounting

Let L_j, M_j, U_j be H5 total-POC lower, median, and upper forecasts. The registered and diagnostic positions are

$$d_j^{\text{int}} = \mathbb{I}\{L_j > 0\}, \qquad d_j^{\text{LF}} = \mathbb{I}\{M_j > 0\}, \qquad d_j^{\text{LS}} = \text{sign}(M_j). \tag{7.6}$$

The latter two were examined after the registered rule abstained and therefore support diagnosis, not confirmation.

For entry P_j^{in}, exit P_j^{out}, position d_j, and one-way proportional cost κ, the documented completed-trade accounting is

$$\begin{aligned} r_j^{\text{gross}} &= d_j \left(\frac{P_j^{\text{out}}}{P_j^{\text{in}}} - 1 \right), \\ r_j^{\text{net}} &= r_j^{\text{gross}} - 2\kappa \, \mathbb{I}\{d_j \neq 0\}, \\ R_{\text{comp}} &= \prod_{j=1}^{N} (1 + r_j^{\text{net}}) - 1. \end{aligned} \tag{7.7}$$

Break-even cost

If mean gross edge is e_{rt} basis points per completed round trip and cost is symmetric, break-even one-way cost is

$$\kappa_{\text{BE}} = \frac{e_{\text{rt}}}{2}. \tag{7.8}$$

For the long-flat diagnostic, $e_{\text{rt}} = 0.4453$ bps, hence $\kappa_{\text{BE}} = 0.2227$ bps per side. The 12-bps base assumption is $12/0.2227 = 53.89$ times break-even. Equivalently, the 24-bps round-trip cost divided by 0.4453 bps of gross edge is 53.89. This arithmetic makes the tested turnover economically infeasible under the base cost.

7.3 Dependence and descriptive block bootstrap

At one-minute spacing, adjacent examples share

$$\frac{63}{64} = 98.4375\% \tag{7.9}$$

of bar rows. Complete five-horizon label windows share

$$\frac{4}{5} = 80\% \tag{7.10}$$

of their minute intervals. A generic autocorrelation effective-sample-size heuristic is

$$n_{\text{eff}} \approx \frac{n}{1 + 2\sum_{\ell=1}^{L} \widehat{\rho}(\ell)}, \tag{7.11}$$

under a stated truncation rule. The recorded H5 heuristic gave an effective fraction near 22.3%; this is diagnostic, not an inferential correction.

The paired daily-block bootstrap sampled four complete UTC-day blocks with replacement, applied the same sampled blocks to model and baseline, and recomputed

$$\Delta^* = \widehat{\mathcal{L}}^*_{\text{base}} - \widehat{\mathcal{L}}^*_{\text{model}}. \tag{7.12}$$

With only four distinct days, its percentile distribution is descriptive and cannot support a confirmatory p-value or broad market inference.

8 Results

8.1 Aggregate probabilistic score

On the development-forward partition, $\mathcal{L}_{\text{model}} = 0.2513909348275368$ and $\mathcal{L}_{\text{base}} = 0.25719276632184623$. Substitution into Equation (7.2) gives

$$100\frac{0.2571927663 - 0.2513909348}{0.2571927663} = 2.2558299665\%. \tag{8.1}$$

This answers the distributional question positively within the bounded development design.

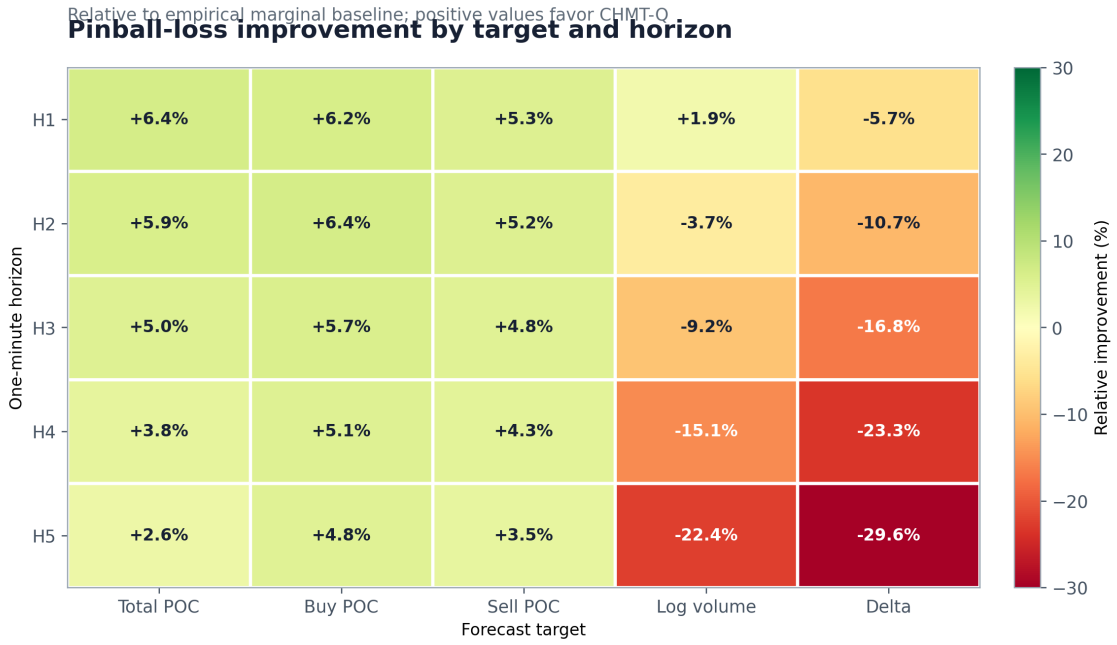


Figure 4: Relative pinball-loss improvement by horizon and target. Positive values favor CHMT-Q. Aggregate weighting masks poor log-volume and delta cells.

All three POC targets improved in every horizon, with diminishing magnitude. Log-volume residual improved only at horizon one and worsened thereafter. Delta pinball loss was worse at every horizon. At H5, target-level relative changes were +2.62%, +4.84%, +3.53%, −22.43%, and −29.57%.

The arithmetic improvement is real on this partition. It is not evidence that a Transformer is necessary, that calibration transfers across regimes, or that the result is statistically significant across independent market periods.

8.2 Five-minute conditional signal

H5 total-POC median MAE was effectively tied with the marginal baseline. Pearson correlation was 0.0774, Spearman correlation 0.0753, and direct-sign MCC 0.0481. A separate anchor-to-future-close direction proxy had 52.67% accuracy, 52.36% balanced accuracy, MCC 0.0475, and Spearman association 0.0780.

Forward endpoint	CHMT-Q	Baseline	Relative result
Weighted normalized pinball	0.251391	0.257193	+2.2558%
H5 total-POC median MAE (ticks)	532.841	532.905	+0.0119%
H5 log-volume median MAE	–	–	24.44% worse
H5 delta median MAE	–	–	44.18% worse

Table 5: Primary forecast and selected location results. Dashes indicate that the table reports relative diagnostics without reproducing an unused physical unit.

Predicted medians had only 15.13% of realized standard deviation. The average 10%–90% interval spanned 1,904.7 ticks, approximately USDT 190.47 at the dataset tick size, and empirical coverage was 81.15%. Every interval crossed zero, so $\widehat{Z} = 1$ in Equation (7.5).

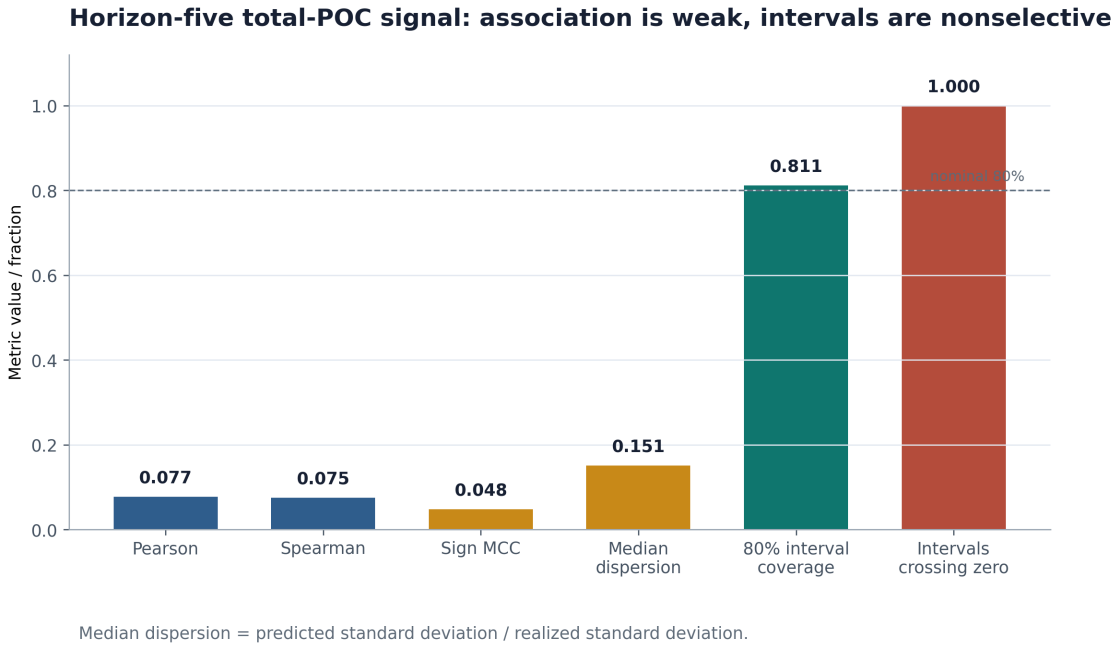


Figure 5: H5 signal diagnostics. Median association and dispersion are weak even though nominal interval coverage is close to 80%.

The coexistence of a shrunk median and broad interval is consistent with learning marginal uncertainty more successfully than conditional location. It does not establish useful directional selection.

8.3 Economic feasibility

The registered interval policy made zero trades because all intervals crossed zero. This is not a software failure: the uncertainty-aware rule abstained. The post hoc median-sign long-flat rule made 484 trades. Its frictionless compounded return was +2.1607%, but its base-cost return was −68.0682%. Mean gross edge was 0.4453 bps per round trip. The diagnostic long-short rule made 1,148 trades, returned +4.989% frictionlessly, and −93.344% at base costs; gross edge was 0.4301 bps.

Policy	Trades	Frictionless	Base-cost	Gross bps / RT
Registered interval long-flat	0	0.00%	0.00%	–
Diagnostic median long-flat	484	+2.1607%	-68.0682%	0.4453
Diagnostic median long-short	1,148	+4.989%	-93.344%	0.4301

Table 6: Economic diagnostics. RT denotes a completed round trip.

The simulator uses minute trade prices and fixed basis-point deductions. It does not reproduce executable quotes, queue position, partial fills, impact, financing, outages, or capacity. It cannot validate live profitability.

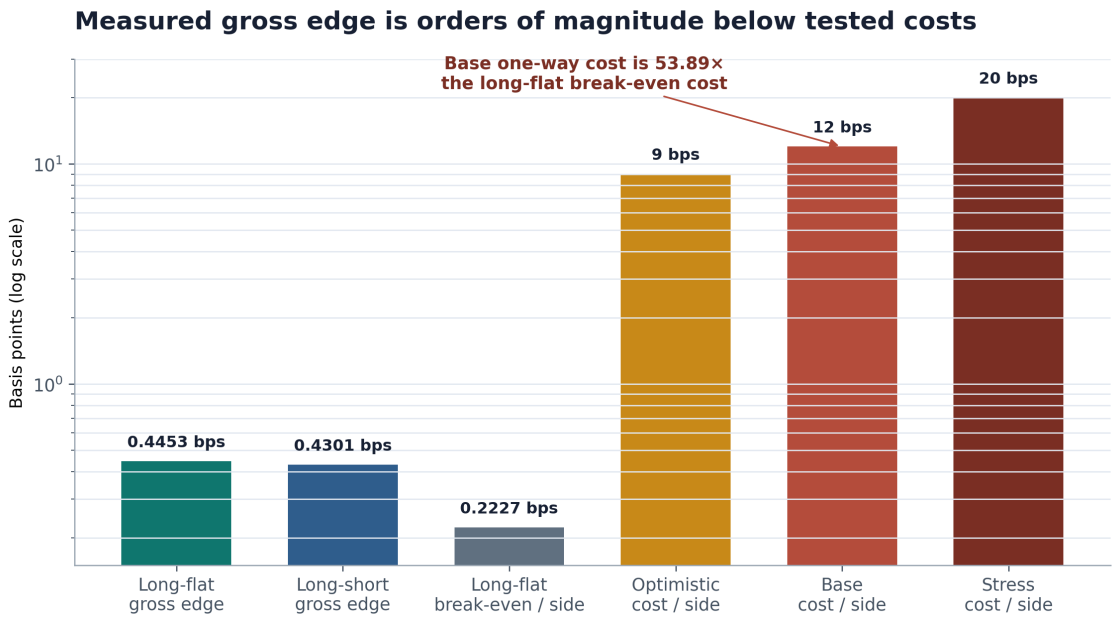


Figure 6: Cost feasibility. Observed gross edge is far below the base round-trip cost of 24 bps.

Because even this simplified accounting is cost-dominated, adding realistic frictions would not rescue the tested policies.

9 Failure-Factor Audit

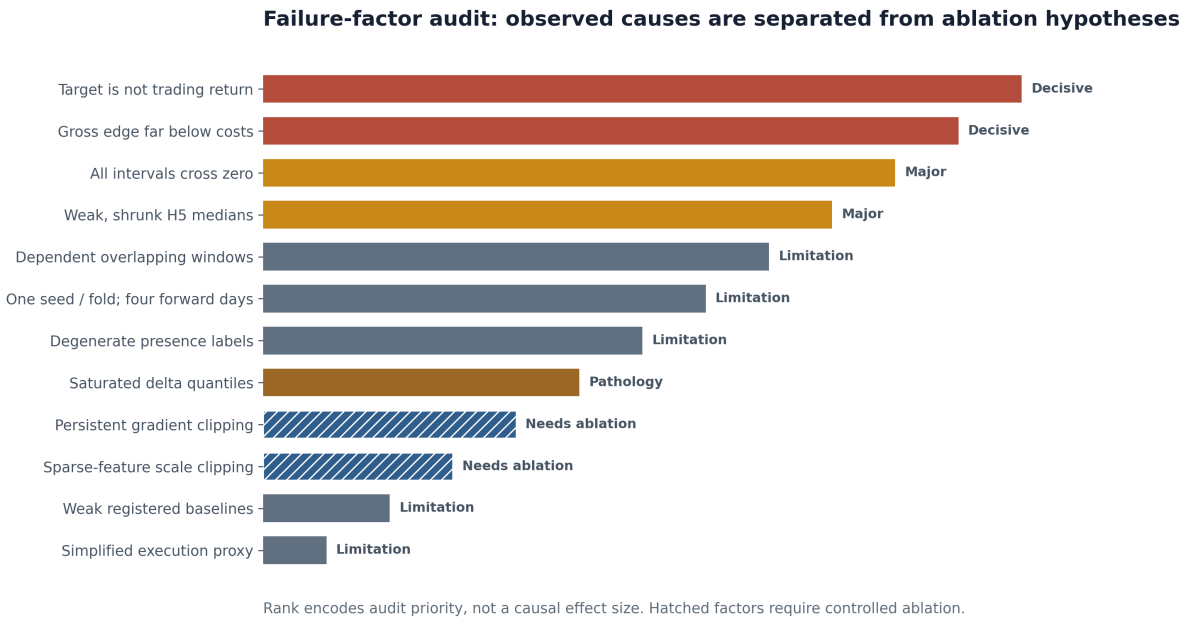


Figure 7: Ranked failure factors. Observed decisive and major factors are separated from plausible contributors requiring prospective ablation.

9.1 Decisive factors

Task-decision mismatch. The model optimized five footprint primitives, not executable return or action. Equation (6.4) shows that only 3.46% of nominal horizon-target weighting concerned H5 total POC. Lower aggregate loss can therefore coexist with an unchanged economic proxy.

Cost dominance. The long-flat gross edge of 0.4453 bps per round trip is structurally incompatible with a 24-bps base round trip. Threshold search can change the selected sample but cannot retroactively turn the observed 53.89-times cost-to-break-even ratio into confirmatory evidence.

Uncertainty without selectivity. Every H5 interval crossed zero. The registered rule correctly abstained. Forcing a median-sign trade discards the uncertainty information that the model was explicitly trained to produce.

9.2 Data, validation, and benchmark limitations

Adjacent examples share most feature and label support. Training comprises 25,758 dependent rows from 18 sessions. With 2.09 million parameters, this is approximately 81 parameters per nominal row and 116,086 per training session. One seed, one fold, two selection days, and four forward days cannot estimate seed instability, regime robustness, venue transfer, or tail behavior. The retired final period removes the planned confirmatory window.

The empirical marginal baseline tests conditional structure relative to a stationary marginal distribution. It does not establish whether attention, hierarchy, or 2.09 million parameters are warranted. Ridge, logistic, boosted-tree, temporal-convolutional, DeepLOB-like, persistence, and seasonal baselines are still required.

9.3 Target and output pathologies

The three POC targets were correlated at roughly 0.93–0.97 and together received

$$\frac{3}{3 + 0.5 + 0.5} = 75\% \tag{9.1}$$

of target weighting. Buy and sell presence were positive for all relevant forward examples, making the auxiliary classifier nearly constant. For an all-positive label, $\mathrm{BCE}(\eta, 1) = \log(1 + e^{-\eta}) \to 0$ as $\eta \to \infty$, so the task offers no directional selectivity.

For delta, the realized median was -0.0345, while the predicted-median distribution had median 0.6695 and bias 0.6881; Pearson association was -0.0107. The derivative

$$\frac{d}{dx} \tanh x = 1 - \tanh^2 x \longrightarrow 0 \quad \text{as} \quad |\tanh x| \longrightarrow 1 \tag{9.2}$$

describes why tail resolution vanishes near the bounds, but does not establish the causal effect of saturation on aggregate performance.

Sparse L1 changes also interacted poorly with robust scaling. When MAD = 0, Equation (5.2) gives $s = 10^{-6}$, so a deviation of only 8×10^{-6} reaches the clipping boundary. Approximately 22.84% of selected mid/bid/ask changes were clipped. Persistent gradient clipping and sparse-feature clipping are observed, but their metric effects require controlled retraining.

Rank	Factor	Evidence status	Implication
1	Target is not trading return	Decisive, observed	Forecast loss is not a PnL objective.
2	Gross edge is far below costs	Decisive, observed	Tested policies are economically infeasible.
3	Every interval crosses zero	Major, observed	Registered policy abstains.
4	Weak, shrunk H5 medians	Major, observed	Aggregate score masks weak location signal.
5	Heavy window overlap	Design limitation	Nominal sample size misleads.
6	One seed/fold; four forward days	Validation limitation	Robustness is unknown.
7	Degenerate presence labels	Wasted capacity	Auxiliary classification adds no selectivity.
8	Saturated delta tails	Output pathology	Bounded uncertainty is uninformative.
9	Persistent gradient clipping	Plausible contributor	Requires optimizer ablation.
10	Sparse-feature clipping	Representation loss	Requires scaler ablation.
11	Weak baseline set	Benchmark limitation	Transformer necessity is untested.
12	Simplified execution	Fidelity limitation	Live performance is unestablished.

Table 7: Failure-factor hierarchy.

10 Threats to Validity

10.1 Internal and statistical conclusion validity

The checkpoint and baseline were frozen before the forward evaluation, but diagnostic policies and several decompositions were developed after inspecting results. They are labeled post hoc. A single run cannot separate architecture effects from initialization. Rows and labels are dependent, and only four calendar days contribute to the estimate. Row-wise confidence intervals would be anti-conservative; four daily blocks do not support a stable tail probability. No significance claim, Sharpe inference, or probability of profitability is reported.

10.2 Construct and external validity

Footprint POC and volume are not executable returns. Minute open/close trade prices are not guaranteed fills. Fixed-bps costs omit dynamic spread, latency, queue priority, impact, and partial execution. The evidence covers one venue, one instrument, one month, one seed, one fold, and four forward days. It does not establish transfer across assets, derivatives, venues, volatility regimes, fee tiers, or live infrastructure.

10.3 Researcher degrees of freedom

The study has one registered forecast endpoint and one registered interval policy, but the audit examines many cells and diagnostics. These analyses generate hypotheses; they do not upgrade the original result. Future work must freeze targets, baselines, thresholds, costs, folds, seeds, and stopping rules before accessing a new confirmatory period.

11 Corrective Research Protocol

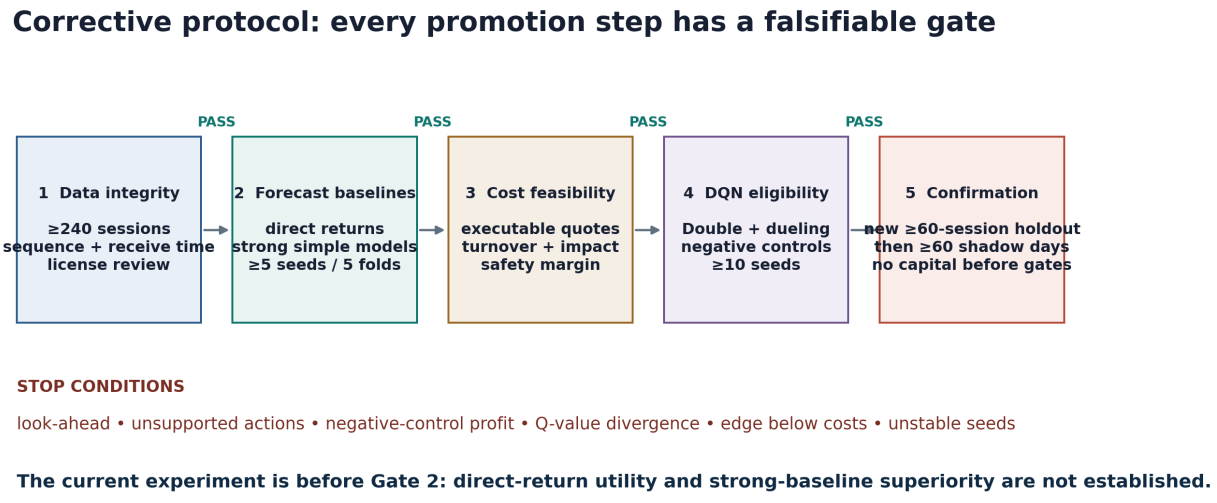


Figure 8: Registered next-study gates. Forecast, execution, and reinforcement learning require separate evidence and a newly sealed period.

11.1 Data and validation gates

For cutoff c_i, maximum lookback L_{lookback}, forecast span $H\Delta$, and label-completeness grace g_{label}, the full dependency interval is

$$J_i = [c_i - L_{\text{lookback}}, c_i + H\Delta + g_{\text{label}}]. \tag{11.1}$$

Prospective train/test assignment must purge any observation whose J_i crosses a fold boundary and apply an embargo at least as wide as the maximum remaining dependency.

The next experiment should acquire at least 240 UTC sessions with exchange event time and, where feasible, local receive time, sequence/checksum integrity, historical instrument metadata, and a documented license review. It should use at least five anchored rolling folds, each with at least 120 training, 20 selection, and 20 disjoint test sessions. At least five Transformer seeds and ten finalist decision-model seeds are required. A new access-controlled holdout of at least 60 sessions must remain target-inaccessible until code, hashes, thresholds, costs, and endpoints are locked, followed by at least 60 shadow-trading days before any capital decision.

11.2 Targets, baselines, and economic gate

Direct executable-return targets should be primary. Footprint targets should remain only if incremental value survives ablation. Redundant POC heads should be selected or orthogonalized; degenerate presence labels removed; and bounded delta modeled with a distribution or transform compatible with its support. Every complex candidate must be compared with frozen persistence/seasonal, ridge/logistic, gradient-boosted, temporal-convolutional, and DeepLOB-like baselines.

The primary endpoint should be net return under a preregistered execution model, reported with turnover, drawdown, exposure, and capacity. The feasibility gate is not positive frictionless return: a lower uncertainty bound for gross edge must exceed estimated all-in cost by a frozen safety margin across folds and seeds.

12 Proposed Transformer–Double-DQN Study

This section is a mathematical proposal only. No DQN has been implemented, trained, backtested, or integrated, and no result in this paper comes from reinforcement learning.

12.1 State, action, transition, and reward

The proposed state combines CHMT-Q embeddings, current target position, realized volatility, spread/liquidity proxies, time-of-day, and risk-budget variables. Actions specify next target position

$$a_t \in \{-1, 0, +1\}, \tag{12.1}$$

rather than ambiguous buy/sell/hold events. With position $p_t = a_t$, mid-price change Δm_{t+1}, transaction-cost function C_t, and bounded risk penalty $\mathcal{R}_t$, shaped training reward would be

$$r_t = p_t \Delta m_{t+1} - C_t(|p_t - p_{t-1}|, \text{spread}_t, \text{liquidity}_t) - \lambda_{\text{risk}} \mathcal{R}_t. \tag{12.2}$$

The reported endpoint must remain raw net PnL under the frozen simulator; shaping is not business value. Historical transitions are valid only under the explicit small-notional assumption that candidate actions do not alter future market observations.

12.2 Double-DQN and dueling derivation

For discount γ, online parameters θ, and target parameters $\bar{\theta}$, Double DQN separates action selection from evaluation:

$$\begin{aligned} a^*_{t+1} &= \arg\max_{a'} Q_\theta(s_{t+1}, a'), \\ y_t^{\text{DDQN}} &= r_t + \gamma(1 - d_t) Q_{\bar{\theta}}(s_{t+1}, a^*_{t+1}), \\ \mathcal{L}_{\text{TD}}(\theta) &= \mathbb{E}_{\mathcal{B}}\Big[H_\delta\Big(Q_\theta(s_t, a_t) - y_t^{\text{DDQN}}\Big)\Big], \end{aligned} \tag{12.3}$$

where d_t marks termination and $\mathcal{B}$ is replay data. The target network would be updated slowly or periodically, with cadence frozen before evaluation.

A dueling head decomposes

$$Q_\theta(s, a) = V_\theta(s) + A_\theta(s, a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A_\theta(s, a'), \tag{12.4}$$

where mean-centering identifies value and advantage. For offline data, a conservative penalty may be ablated:

$$\mathcal{L}_{\text{CQL}} = \alpha\, \mathbb{E}_s \left[\log \sum_a e^{Q_\theta(s,a)} - \mathbb{E}_{a \sim \mathcal{D}(\cdot|s)} Q_\theta(s, a) \right]. \tag{12.5}$$

This discourages unsupported actions from receiving arbitrarily high values; it does not validate the simulator or create action support.

12.3 Eligibility and stopping gates

Required controls include behavior cloning, no-change policies, randomized rewards, shifted features, and permuted action labels. Profitable negative controls indicate leakage or simulator error. Training must stop if action support is inadequate, Q-values diverge, controls profit, or seed dispersion exceeds a preregistered tolerance. Discount, target-update cadence, replay sampling, and multi-step return length remain unspecified and must be frozen before implementation.

Gate	Minimum condition	Failure consequence
Data integrity	Sequence/receive-time and no-lookahead tests pass	Stop
Forecast utility	Direct-return model beats frozen simple baselines across folds	Do not train DQN
Cost feasibility	Gross edge clears all-in costs with safety margin	Do not train DQN
Offline support	Candidate actions have adequate state support	Restrict actions or stop
Robustness	At least 5 encoder and 10 finalist DQN seeds	No promotion
Confirmation	New, sealed, at least 60-session holdout	No performance claim
Operations	At least 60 live shadow days	No capital allocation

Table 8: Prerequisites for any future DQN study.

13 Reproducibility, Availability, and Ethics

13.1 Artifact ledger

Table 9: Immutable local artifact ledger.

Artifact	SHA-256 or identifier	Role
Compiled dataset	`565c1d8e912ee426a1a92572d5be481ec5168d54d1ae53bd77fdd70d888e4df6`	Dataset identity
Dataset manifest	`68ec9eca21d2a30546435738e8e07d565a64edbf01825246db659a4c1bb98046`	Source/shard manifest
Compiler report	`cc3f9208e2410b41ca9836ef762c15cfd91d3b9c704c2d198c438425f4c9e8ba`	Compilation audit
Dependency lock	`48792cf2789a576553fc29b1b727e97342d28fed1fdbeb91c31135a84bfff898`	Environment identity
Source split	`7f5254f62a3d70a6f02e41b1991f25d3bd8338e7aa9b25c5fa19737f70a7bb8e`	Partition identity
Input/target scaler	`29a733fac9986eb7e879555b6813f8d35b06c98dee9041057e6fc4287bd50880`	Scaling state
Training run	`25be2f446a75ad37120bba956ea86e46c0ecefa004ab3343d6745cdbee102fb3`	Experiment identity
Study protocol	`885fcd0fee23115a4a5480ec94d95f436e251ec0caf796f9686867ed7e72f5d0`	Frozen protocol
Prediction tape	`27d9ffbe41b18868b8a45bf9c48c936908de7394dd6ecfec7062cc531cc7e1f0`	Model outputs
Study results	`083da170146585cf1a16dde56e2da02697194640ed7b053d3298d8e3775bee6a`	Results ledger
Failure diagnostics	`ece5b18f8a8455f7edac8491b6d23362fa14f618129b1193b1febf14bc89505e`	Audit results
DQN plan	`d5e81f203619c55754c940df7e92ec5995e4db41b3a7d6f3db4137318f7a502c`	Future-work specification

The implementation and local artifacts are private research materials. Raw exchange archives, derived Arrow data, the prediction tape, weights, and checkpoints are not redistributed. Reproduction requires independent data acquisition under applicable OKX terms and regeneration under the disclosed tensor, split, training, and evaluation contracts. Hashes establish local identity; they do not substitute for public data availability.

13.2 Risk and declarations

The model is unsuitable for capital allocation. Any later live test requires shadow mode, hard position/loss limits, kill switches, human oversight, and monitoring for drift, latency, execution quality, and turnover.

Funding To be completed by the authors.
Competing interests
To be completed by the authors.
Author contributions
To be completed using the selected venue taxonomy.
Generative tools To be declared according to venue policy. Human authors must independently verify every claim, equation, reference, table, and figure.
Data and code Private local research artifacts; no redistribution claim.
Human or animal participants
Not applicable.

14 Conclusion

CHMT-Q improved a registered aggregate proper score by 2.26% on a four-session development-forward partition. That result did not imply a useful trading signal. H5 medians were weak and under-dispersed, every uncertainty interval crossed zero, the registered policy made no trades, and diagnostic policies lost heavily once fixed costs were applied. The observed long-flat gross edge was 53.89 times too small for the base cost assumption.

The strongest contribution is methodological. Proper scores, conditional signal, and executable economics are distinct evidentiary levels. A model should not advance because it wins at the first while failing or remaining untested at the others. The current checkpoint remains research-only and non-promotable. A larger, access-controlled, multi-fold, multi-seed study with direct return targets, strong baselines, realistic execution, and a new sealed holdout is required before reconsideration. The proposed Transformer–Double-DQN design is one testable option inside that protocol, not a remedy already demonstrated.

A Detailed Derivations and Checks

A.1 Ordered-head proof after bounded transformation

Lemma 1. *Let* $q_L = c - \mathrm{softplus}(a)$, $q_M = c$, $q_U = c + \mathrm{softplus}(b)$. *If* $f : \mathbb{R} \to \mathbb{R}$ *is nondecreasing, then* $f(q_L) \le f(q_M) \le f(q_U)$.

Proof. $\mathrm{softplus}(a), \mathrm{softplus}(b) > 0$ implies $q_L < q_M < q_U$. The definition of a nondecreasing function preserves both inequalities. Taking $f = \tanh$ proves ordered bounded delta quantiles. Strict mathematical order may become numerical equality in saturated floating-point tails. □

A.2 Aggregate score decomposition

Let N_{hk} denote the weighted sum of pinball losses for cell (h, k) and D_{hk} its mask/weight denominator. Then

$$\mathcal{L}_Q = \frac{\sum_{h,k} N_{hk}}{\sum_{h,k} D_{hk}} = \sum_{h,k} \omega_{hk} \mathcal{L}_{hk}, \qquad \omega_{hk} = \frac{D_{hk}}{\sum_{h',k'} D_{h'k'}}. \tag{A.1}$$

An aggregate improvement therefore permits negative cell improvements whenever positive, high-weight cells dominate. In this run, the three correlated POC targets carry most target weight, while log-volume and delta deterioration is downweighted.

A.3 Proper-score and trading-value counterexample

Consider $X \sim \mathcal{N}(0, 1)$, $Y = X + \epsilon$ with $\epsilon \sim \mathcal{N}(0, \sigma^2)$, and executable return $R = \eta$ with η independent and symmetric. The conditional median forecast $q_{0.5}(X) = X$ strictly improves expected absolute/pinball loss over the unconditional median 0 whenever X is nondegenerate. Yet $\mathbb{E}[R \mid X, Y] = 0$, so for any position a,

$$\mathbb{E}[aR - C|a| \mid X, Y] = -C|a| \le 0, \tag{A.2}$$

with strict loss for $a \neq 0$ and $C > 0$. This formalizes why score improvement cannot, without a target-decision bridge, establish edge.

A.4 Cost compounding

For constant small gross arithmetic edge g and round-trip cost $2c$ over N trades,

$$R_N = (1 + g - 2c)^N - 1. \tag{A.3}$$

The first-order log approximation is

$$\log(1 + R_N) = N \log(1 + g - 2c) \approx N(g - 2c) - \frac{N}{2}(g - 2c)^2. \tag{A.4}$$

Even a positive frictionless g compounds rapidly downward when $2c > g$. This explains the large negative cost-adjusted returns without invoking leverage.

A.5 Parameter ledger reconciliation

For a biased linear map $a \to b$,

$$P_{\text{linear}}(a, b) = ab + b, \tag{A.5}$$

and a biased SwiGLU block $d \to f \to d$ has

$$P_{\text{SwiGLU}}(d, f) = 3df + 2f + d. \tag{A.6}$$

Together with attention projections, normalizers, LayerScale vectors, embeddings, gates, and learned head temperatures, the component ledger is:

Component	**Trainable parameters**
Bar, L1, and level projections	4,768
Quality MLP	18,688
Bar time MLP	4,768
L1 patch-time MLP	4,832
L2 patch-time MLP	4,832
Modality embeddings	384
L2 side/rank embeddings	384
L2 level CHMT block	10,660
Book query, norm, and 32-to-128 projection	4,288
L1/L2 patch embeddings and pool MLPs	9,346
Nine temporal CHMT blocks	1,489,608
Two decoder blocks	463,632
Three modality gate MLPs	49,539
Fusion norm and context projection	16,640
Horizon queries and horizon MLP	4,992
Shared 128-to-17 head	2,193
Total	**2,089,554**

Table 10: Exact implemented parameter ledger. The total includes 36 learned per-head Q/K temperatures omitted from the superseded 2,089,518 arithmetic.

B Notation

Table 11: Principal notation.

Symbol	**Meaning**
B	Batch size
c	Forecast decision cutoff
$\mathcal{F}_c$	Information available by cutoff
H	Number of direct forecast horizons, five
p_{ref}	Causal reference price
δ_p	Instrument tick size

Symbol	Meaning
ℓ_{ref}	Causal trailing log-volume reference
d, d_h, d_{ff}	Model, attention-head, and feed-forward widths
m	Boolean validity/eligibility mask
α_m	Normalized modality gate
Q_τ	Conditional quantile at level τ
ρ_τ	Pinball loss
θ	Learned model parameters
p_t	Target trading position
κ_{BE}	Break-even one-way transaction cost

C Registered and Post Hoc Analyses

Analysis	Status	Interpretation
Weighted normalized pinball versus empirical marginal	Registered forecast endpoint	Development evidence
Interval long-flat policy at base costs	Registered economic endpoint	Zero trades; development evidence
Median-sign long-flat and long-short policies	Diagnostic post hoc	Cost-feasibility diagnosis only
Anchor-close candle direction proxy	Diagnostic post hoc	Approximate downstream association
Per-target/per-horizon heatmaps	Diagnostic decomposition	Explains aggregate heterogeneity
Recent-three-day and hour baselines	Post hoc; not in primary comparison	Hypothesis generation only
DQN architecture and validation plan	Proposed future work	No performance evidence

Table 12: Status of analyses.

D Submission Completion Items

1. Complete affiliation, correspondence, ORCID, funding, competing-interest, contribution, and tool-use fields as required.
2. Select the target venue and apply its exact class and bibliography requirements without altering the empirical claims.
3. Independently recompute and verify every quantitative table and figure from immutable artifacts.
4. Preserve the holdout-integrity disclosure; never relabel June 25–30 as sealed, untouched, or confirmatory.
5. Confirm that the venue permits private code/data with a regeneration contract, or choose an appropriate controlled archive and release policy.

References

[1] A. Vaswani et al. Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 2017. https://arxiv.org/abs/1706.03762.

[2] P. Shaw, J. Uszkoreit, and A. Vaswani. Self-attention with relative position representations. *NAACL-HLT*, pages 464–468, 2018. https://doi.org/10.18653/v1/N18-2074.

[3] B. Zhang and R. Sennrich. Root mean square layer normalization. *Advances in Neural Information Processing Systems*, 32, 2019. https://arxiv.org/abs/1910.07467.

[4] N. Shazeer. GLU variants improve Transformer. Technical report, 2020. https://arxiv.org/abs/2002.05202.

[5] Y. Nie et al. A time series is worth 64 words: Long-term forecasting with Transformers. *International Conference on Learning Representations*, 2023. https://arxiv.org/abs/2211.14730.

[6] Y. Liu et al. iTransformer: Inverted Transformers are effective for time series forecasting. *International Conference on Learning Representations*, 2024. https://arxiv.org/abs/2310.06625.

[7] B. Lim, S. O. Arik, N. Loeff, and T. Pfister. Temporal fusion Transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4):1748–1764, 2021. https://doi.org/10.1016/j.ijforecast.2021.03.012.

[8] A. Zeng et al. Are Transformers effective for time series forecasting? *AAAI*, 37(9):11121–11128, 2023. https://doi.org/10.1609/aaai.v37i9.26317.

[9] Z. Zhang, S. Zohren, and S. Roberts. DeepLOB: Deep convolutional neural networks for limit order books. *IEEE Transactions on Signal Processing*, 67(11):3001–3012, 2019. https://doi.org/10.1109/TSP.2019.2907260.

[10] J. Wallbridge. Transformers for limit order books. Technical report, 2020. https://arxiv.org/abs/2003.00130.

[11] Z. Zhang and S. Zohren. Multi-horizon forecasting for limit order books: Novel deep learning approaches and hardware acceleration using intelligent processing units. Technical report, 2021. https://arxiv.org/abs/2105.10430.

[12] A. Ntakaris et al. Benchmark dataset for mid-price forecasting of limit order book data with machine learning methods. *Journal of Forecasting*, 37(8):852–866, 2018. https://doi.org/10.1002/for.2543.

[13] P. N. Kolm, J. Turiel, and N. Westray. Deep order flow imbalance: Extracting alpha at multiple horizons from the limit order book. *Mathematical Finance*, 33(4):1044–1081, 2023. https://doi.org/10.1111/mafi.12413.

[14] L. Lucchese, M. S. Pakkanen, and A. E. D. Veraart. The short-term predictability of returns in order book markets: A deep learning perspective. *International Journal of Forecasting*, 40(4):1607–1626, 2024. https://doi.org/10.1016/j.ijforecast.2024.02.001.

[15] A. Briola, S. Bartolucci, and T. Aste. Deep limit order book forecasting: A microstructural guide. *Quantitative Finance*, 2025. https://doi.org/10.1080/14697688.2025.2522911.

[16] T. Gneiting and A. E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007. https://doi.org/10.1198/016214506000001437.

[17] Y. Romano, E. Patterson, and E. J. Candes. Conformalized quantile regression. *Advances in Neural Information Processing Systems*, 32, 2019. https://arxiv.org/abs/1905.03222.

[18] J. Sirignano and R. Cont. Universal features of price formation in financial markets: Perspectives from deep learning. *Quantitative Finance*, 19(9):1449–1459, 2019. https://doi.org/10.1080/14697688.2019.1622295.

[19] H. White. A reality check for data snooping. *Econometrica*, 68(5):1097–1126, 2000. https://doi.org/10.1111/1468-0262.00152.

[20] D. N. Politis and J. P. Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89(428):1303–1313, 1994. https://doi.org/10.1080/01621459.1994.10476870.

[21] C. R. Harvey, Y. Liu, and H. Zhu. . . . and the cross-section of expected returns. *Review of Financial Studies*, 29(1):5–68, 2016. https://doi.org/10.1093/rfs/hhv059.

[22] V. Mnih et al. Human-level control through deep reinforcement learning. *Nature*, 518:529–533, 2015. https://doi.org/10.1038/nature14236.

[23] H. van Hasselt, A. Guez, and D. Silver. Deep reinforcement learning with Double Q-learning. *AAAI*, 30(1):2094–2100, 2016. https://doi.org/10.1609/aaai.v30i1.10295.

[24] Z. Wang, T. Schaul, M. Hessel, H. van Hasselt, M. Lanctot, and N. de Freitas. Dueling network architectures for deep reinforcement learning. *Proceedings of Machine Learning Research*, 48:1995–2003, 2016. https://proceedings.mlr.press/v48/wangf16.html.

[25] W. Dabney, M. Rowland, M. Bellemare, and R. Munos. Distributional reinforcement learning with quantile regression. *AAAI*, 32(1), 2018. https://doi.org/10.1609/aaai.v32i1.11791.

[26] A. Kumar, A. Zhou, G. Tucker, and S. Levine. Conservative Q-learning for offline reinforcement learning. *Advances in Neural Information Processing Systems*, 33:1179–1191, 2020. https://arxiv.org/abs/2006.04779.

[27] S. Levine, A. Kumar, G. Tucker, and J. Fu. Offline reinforcement learning: Tutorial, review, and perspectives on open problems. Technical report, 2020. https://arxiv.org/abs/2005.01643.

[28] P. Henderson et al. Deep reinforcement learning that matters. *AAAI*, 32(1), 2018. https://doi.org/10.1609/aaai.v32i1.11694.

[29] D. H. Bailey, J. M. Borwein, M. Lopez de Prado, and Q. J. Zhu. The probability of backtest overfitting. *Journal of Computational Finance*, 20(4):39–69, 2017. https://doi.org/10.21314/JCF.2016.322.

[30] J. Racine. Consistent cross-validatory model-selection for dependent data: hv-block cross-validation. *Journal of Econometrics*, 99(1):39–61, 2000. https://doi.org/10.1016/S0304-4076(00)00030-0.